

Jingxuan (Jasper) Wu

✉ jingxwu@unc.edu [🏠 johnny221b.github.io](https://github.com/johnny221b) [🎓 Google Scholar](#)

Last updated: May, 2026

Research Interest

My research aims to build controllable generative systems through the integration of diffusion/flow-based models and multi-agent systems: (1) **Generative Models**—enabling diverse, fine-grained controllable generation under limited supervision through both training-free inference-time control and RL-based training; (2) **Agentic Systems**—designing multi-agent systems that interact, coordinate, and adapt through feedback.

Education

The University of North Carolina at Chapel Hill
M.S. in Statistics

Aug. 2025 – Dec. 2026 (Expected)
Chapel Hill, NC, USA

The Chinese University of Hong Kong, Shenzhen
B.S. (Hons) in Data Science and Big Data Technology

Sep. 2021 – May 2025
Shenzhen, China

Selected Publications

(* indicates equal contribution)

Conference & Workshop Papers

- [C1] **Jingxuan Wu**, Zhenglin Wan, Xingrui Yu, Yuzhe Yang, Bo An, and Ivor Tsang. **“OSCAR: Orthogonal Stochastic Control for Alignment-Respecting Diversity in Flow Matching”**. 2025.
ICML 2026.
- [C2] Zhenglin Wan, **Jingxuan Wu**, Xingrui Yu, Chubin Zhang, Mingcong Lei, Bo An, and Ivor Tsang. **“FM-IRL: Flow-Matching for Reward Modeling and Policy Regularization in Reinforcement Learning”**. 2025.
ICML 2026.
- [C3] **Jingxuan Wu**, Zhenglin Wan, Xingrui Yu, Yuzhe Yang, Yiqiao Huang, Ivor Tsang, and Yang You. **“Time-Annealed Perturbation Sampling: Diverse Generation for Diffusion Language Models”**. 2026.
ICML 2026 Workshop on SPIGM.

Journal Papers

- [J1] Chi Li, Jianfeng Mao, Lingyi Li, **Wu, Jingxuan**, Lianmin Zhang, Jianyu Zhu, and Zibin Pan. **“Flight delay propagation modeling: Data, Methods, and Future opportunities”**. In: *Transportation Research Part E: Logistics and Transportation Review* 185 (2024), p. 103525.

Preprints & Technical Reports

- [P1] Wang, Hanzhao*, **Wu, Jingxuan***, Pan, Yu, Li, Yumeng, Wang, Yansong, Liu, Helang, Wang, Fuqiang, and Chen, Guanting. **“LLM-Powered Predictive Decision-Making for Sustainable Data Center Operations”**. OpenReview. 2024.
In Preparation for *Manufacturing & Service Operations Management*.

Research Experience

Statistics & Operations Research, University of North Carolina, Chapel Hill
Advisors: [Prof. Guanting Chen](#), [Prof. Xiaocheng Li](#), [Prof. Wenjia Ba](#)

Jan. 2024 – Present
Chapel Hill, NC, USA

- **Adaptive Preference Learning for Generative Design via Hierarchical Proxy-Human Agents:** Built an end-to-end experimental pipeline for Diffusion Language Models, integrating backbones, inference-time interventions, decoding baselines, and automatic evaluation of semantic diversity and generation quality.
- **Resource Allocation Optimization for GPUs using LLMs:** Developed an LLM-driven GPU scheduling pipeline that predicts runtime/energy from code embeddings, models arrivals with a non-homogeneous Poisson process, and evaluates an online allocator on datacenter traces to improve performance.

A*STAR Centre for Frontier AI Research

May. 2025 – Present

Advisors: [Prof. Xingrui Yu](#), [Prof. Ivor Tsang](#), [Prof. Bo An](#)

Singapore

- **Time-Annealed Perturbation Sampling for Diffusion Language Models (TAPS):** Proposed a training-free, scalable inference-time perturbation scheme with time-dependent context annealing and quality-preserving rescaling to boost generation diversity in Diffusion Language Models.
- **Guidance Mechanisms to Enhance Diversity in Flow Matching:** Developed a training-free inference-time control framework (OSCAR) that improves set-level diversity while preserving quality/alignment via feature-space volume maximization and orthogonal stochastic control injected into the base flow.
- **Flow-Matching Rewards & Policy Regularization:** Introduced a teacher–student FM-IRL framework where a flow-matching teacher provides reward shaping and policy regularization, enabling stable online updates and improved efficiency/robustness under limited or suboptimal demonstrations.

School of Data Science, CUHK-Shenzhen

Dec. 2023 – Nov. 2024

Advisors: [Prof. Jianfeng Mao](#)

Shenzhen, China

- **Research on Network-wide Delay Propagation Dynamics Prediction:** Modeled network-wide flight delay propagation with a heterogeneous time-varying SIS process and an adaptive graph learning framework (GAT-based) to learn infection/recovery dynamics, achieving improved multi-horizon delay-status forecasting.
- **Study on Flight Delay Propagation Modeling:** Conducted a structured review of 40+ studies on airport delay propagation across statistical/econometric/queueing paradigms, synthesizing assumptions, strengths/limitations, and data-scale considerations into a comprehensive survey manuscript.

Projects

CaveAgent | Python, LLM, AST, LiteLLM, Asyncio, Security Validation

Oct. 2025 – Jan. 2026

🔗🔗 *A tool-augmented agent framework enabling stateful runtime management and code-based function calling.*

- Contributed to the development of the persistent Python runtime modules, supporting seamless object injection and retrieval across multi-turn LLM interactions.
- Implemented specific AST-based security rules to validate and sanitize LLM-generated code before execution.
- Developed discovery and registry mechanisms for the Agent Skills open standard, allowing modular injection of functions and types into the runtime.
- Participated in benchmarking and performance evaluation across multiple LLM backends (OpenAI, Claude, Gemini) to validate framework robustness.

References available upon request.